

Predicting Gender from OkCupid Profiles

By *Sergio Avalos S.*

1. Introduction

In online dating platforms, user demographics play a crucial role in improving matchmaking algorithms and personalizing user experiences. One key demographic attribute is gender, which influences profile visibility, recommendations, and overall platform engagement. This project aims to build a machine learning model that predicts a user's gender based on various personal attributes from their OkCupid profile. By leveraging structured data such as age, body type, and lifestyle habits, this project seeks to determine how accurately gender can be inferred from available profile information.

1.1. Scoping

Project Goals

The primary research question that will be answered is whether an OkCupid user's gender can be predicted using other variables from their profiles. This project is important since understanding user demographics can help improve recommendations and personalize user experiences.

Data

The project has one dataset provided by Codecademy called `profiles.csv`. In the data, each row represents an OkCupid user, and the columns contain responses to their user profiles, including multiple-choice and short-answer questions. The target variable for prediction is `sex`, which can take values "m" or "f".

The features selected for this analysis include:

- **Age**
- **Body Type**
- **Diet, Drinks, Smokes, Drugs**
- **Education & Job**
- **Height**
- **Religion**
- **Orientation**

Some columns have missing values, and appropriate preprocessing techniques will be used to handle them.

Analysis

This solution will use descriptive statistics and data visualization to explore the distribution, count, and relationships between variables. Since the goal of the project is to make predictions on the user's gender, classification algorithms from the supervised learning family of machine learning models will be implemented. Potential models to explore include:

- Logistic Regression
- Decision Trees
- Random Forest
- Support Vector Machines

Feature engineering, encoding categorical variables, and handling missing data will also be key steps in preparing the dataset for modeling.

Evaluation

The project will conclude with the evaluation of the machine learning model selected using a validation dataset. The output of the predictions can be assessed using a confusion matrix, and metrics such as accuracy, precision, recall, and F1 scores. The final model performance will determine how well gender can be predicted based on the chosen features.

2. Data Preparation and Exploration

2.1. Importing the Modules

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns

df = pd.read_csv('profiles.csv', encoding='utf-8')
print(df.head())
```

	age	body_type	diet	drinks	drugs \
0	22	a little extra	strictly anything	socially	never
1	35	average	mostly other	often	sometimes
2	38	thin	anything	socially	NaN
3	23	thin	vegetarian	socially	NaN
4	29	athletic	NaN	socially	never

	education \
0	working on college/university
1	working on space camp
2	graduated from masters program
3	working on college/university
4	graduated from college/university

	essay0 \
0	about me: \n \ni would love to think...
1	i am a chef: this is what that means. \n1...
2	i'm not ashamed of much, but writing public te...
3	i work in a library and go to school. . .
4	hey how's it going? currently vague on the pro...

	essay1 \
0	currently working as an international agent fo...
1	dedicating everyday to being an unbelievable b...
2	i make nerdy software for musicians, artists, ...
3	reading things written by old dead people
4	work work work work + play

	essay2 \
0	making people laugh. \nranting about a go...
1	being silly. having ridiculous amonts of fun w...
2	improvising in different contexts. alternating...
3	playing synthesizers and organizing books acco...
4	creating imagery to look at: \nhttp://bag...

	essay3 ... \
0	the way i look. i am a six foot half asian, ha... ...
1	NaN ...
2	my large jaw and large glasses are the physica... ...
3	socially awkward but i do my best ...
4	i smile a lot and my inquisitive nature ...

	location \
0	south san francisco, california
1	oakland, california
2	san francisco, california
3	berkeley, california
4	san francisco, california

	offspring orientation \
0	doesn't have kids, but might want them straight
1	doesn't have kids, but might want them straight
2	NaN straight
3	doesn't want kids straight
4	NaN straight

```

                pets                                religion sex \
0 likes dogs and likes cats    agnosticism and very serious about it    m
1 likes dogs and likes cats    agnosticism but not too serious about it    m
2                has cats                                NaN    m
3                likes cats                                NaN    m
4 likes dogs and likes cats                                NaN    m

```

```

                sign    smokes \
0                gemini    sometimes
1                cancer        no
2 pisces but it doesn't matter        no
3                pisces        no
4                aquarius        no

```

```

                speaks    status
0                english    single
1 english (fluently), spanish (poorly), french (...)    single
2                english, french, c++    available
3                english, german (poorly)    single
4                english    single

```

[5 rows x 31 columns]

Data characteristics

The columns in the dataset include:

- **age:** continuous variable of age of user
- **body_type:** categorical variable of body type of user
- **diet:** categorical variable of dietary information
- **drinks:** categorical variable of alcohol consumption
- **drugs:** categorical variable of drug usage
- **education:** categorical variable of educational attainment
- **ethnicity:** categorical variable of ethnic backgrounds
- **height:** continuous variable of height of user
- **income:** continuous variable of income of user
- **job:** categorical variable of employment description
- **offspring:** categorical variable of children status
- **orientation:** categorical variable of sexual orientation
- **pets:** categorical variable of pet preferences
- **religion:** categorical variable of religious background
- **sex:** categorical variable of gender
- **sign:** categorical variable of astrological symbol
- **smokes:** categorical variable of smoking consumption
- **speaks:** categorical variable of language spoken
- **status:** categorical variable of relationship status
- **last_online:** date variable of last login
- **location:** categorical variable of user locations

And a set of open short-answer responses to:

- **essay0:** My self summary
- **essay1:** What I'm doing with my life
- **essay2:** I'm really good at
- **essay3:** The first thing people usually notice about me
- **essay4:** Favorite books, movies, show, music, and food
- **essay5:** The six things I could never do without
- **essay6:** I spend a lot of time thinking about
- **essay7:** On a typical Friday night I am
- **essay8:** The most private thing I am willing to admit
- **essay9:** You should message me if...

For this project's purposes, we will be dropping the essay columns.

```
In [2]: df = df.drop(columns=['essay0', 'essay1', 'essay2', 'essay3', 'essay4', 'essay5', '
print(df.head())
```

	age	body_type	diet	drinks	drugs	\
0	22	a little extra	strictly anything	socially	never	
1	35	average	mostly other	often	sometimes	
2	38	thin	anything	socially	NaN	
3	23	thin	vegetarian	socially	NaN	
4	29	athletic	NaN	socially	never	

	education	ethnicity	height	income	\
0	working on college/university	asian, white	75.0	-1	
1	working on space camp	white	70.0	80000	
2	graduated from masters program	NaN	68.0	-1	
3	working on college/university	white	71.0	20000	
4	graduated from college/university	asian, black, other	66.0	-1	

	job	...	location	\
0	transportation	...	south san francisco, california	
1	hospitality / travel	...	oakland, california	
2	NaN	...	san francisco, california	
3	student	...	berkeley, california	
4	artistic / musical / writer	...	san francisco, california	

	offspring	orientation	\
0	doesn't have kids, but might want them	straight	
1	doesn't have kids, but might want them	straight	
2	NaN	straight	
3	doesn't want kids	straight	
4	NaN	straight	

	pets	religion	sex	\
0	likes dogs and likes cats	agnosticism and very serious about it	m	
1	likes dogs and likes cats	agnosticism but not too serious about it	m	
2	has cats	NaN	m	
3	likes cats	NaN	m	
4	likes dogs and likes cats	NaN	m	

	sign	smokes	\
0	gemini	sometimes	
1	cancer	no	
2	pisces but it doesn't matter	no	
3	pisces	no	
4	aquarius	no	

	speaks	status
0	english	single
1	english (fluently), spanish (poorly), french (...)	single
2	english, french, c++	available
3	english, german (poorly)	single
4	english	single

[5 rows x 21 columns]

2.2. Data Cleaning and Wrangling

We notice that the features `sign`, `religion`, `ethnicity` and `speaks` could be simplified by getting just the first word of the sentence:

```
In [3]: df.sign = df.sign.str.split().str.get(0)
df.religion = df.religion.str.split().str.get(0)
df.ethnicity = df.ethnicity.str.split().str.get(0)
df.ethnicity = df.ethnicity.str.replace('[,]', '', regex=True)
df.speaks = df.speaks.str.split().str.get(0)
df.speaks = df.speaks.str.replace('[,]', '', regex=True)
print(df.head())
```

	age	body_type	diet	drinks	drugs	\
0	22	a little extra	strictly anything	socially	never	
1	35	average	mostly other	often	sometimes	
2	38	thin	anything	socially	NaN	
3	23	thin	vegetarian	socially	NaN	
4	29	athletic	NaN	socially	never	

	education	ethnicity	height	income	\
0	working on college/university	asian	75.0	-1	
1	working on space camp	white	70.0	80000	
2	graduated from masters program	NaN	68.0	-1	
3	working on college/university	white	71.0	20000	
4	graduated from college/university	asian	66.0	-1	

	job	...	location	\
0	transportation	...	south san francisco, california	
1	hospitality / travel	...	oakland, california	
2	NaN	...	san francisco, california	
3	student	...	berkeley, california	
4	artistic / musical / writer	...	san francisco, california	

	offspring	orientation	\
0	doesn't have kids, but might want them	straight	
1	doesn't have kids, but might want them	straight	
2	NaN	straight	
3	doesn't want kids	straight	
4	NaN	straight	

	pets	religion	sex	sign	smokes	speaks	\
0	likes dogs and likes cats	agnosticism	m	gemini	sometimes	english	
1	likes dogs and likes cats	agnosticism	m	cancer	no	english	
2	has cats	NaN	m	pisces	no	english	
3	likes cats	NaN	m	pisces	no	english	
4	likes dogs and likes cats	NaN	m	aquarius	no	english	

	status
0	single
1	single
2	available
3	single
4	single

[5 rows x 21 columns]

Our location feature mentions both city and state, for this project we will only use the state:

```
In [4]: df["location"] = df["location"].apply(lambda x: x.split(", ")[-1] if pd.notnull(x)
print(df.head())
```

	age	body_type	diet	drinks	drugs	\
0	22	a little extra	strictly anything	socially	never	
1	35	average	mostly other	often	sometimes	
2	38	thin	anything	socially	NaN	
3	23	thin	vegetarian	socially	NaN	
4	29	athletic	NaN	socially	never	

	education	ethnicity	height	income	\
0	working on college/university	asian	75.0	-1	
1	working on space camp	white	70.0	80000	
2	graduated from masters program	NaN	68.0	-1	
3	working on college/university	white	71.0	20000	
4	graduated from college/university	asian	66.0	-1	

	job	...	location	\
0	transportation	...	california	
1	hospitality / travel	...	california	
2	NaN	...	california	
3	student	...	california	
4	artistic / musical / writer	...	california	

	offspring	orientation	\
0	doesn't have kids, but might want them	straight	
1	doesn't have kids, but might want them	straight	
2	NaN	straight	
3	doesn't want kids	straight	
4	NaN	straight	

	pets	religion	sex	sign	smokes	speaks	\
0	likes dogs and likes cats	agnosticism	m	gemini	sometimes	english	
1	likes dogs and likes cats	agnosticism	m	cancer	no	english	
2	has cats	NaN	m	pisces	no	english	
3	likes cats	NaN	m	pisces	no	english	
4	likes dogs and likes cats	NaN	m	aquarius	no	english	

	status
0	single
1	single
2	available
3	single
4	single

[5 rows x 21 columns]

For this project, the features `last_online` , `offspring` and `pets` won't be considered neither.

```
In [5]: df = df.drop(columns=['last_online', 'offspring', 'pets'])
print(df.head())
```


	age	body_type	diet	drinks	drugs	\
0	22	a little extra	strictly anything	socially	never	
1	35	average	mostly other	often	sometimes	
2	38	thin	anything	socially	NaN	
3	23	thin	vegetarian	socially	NaN	
4	29	athletic	NaN	socially	never	

	education	ethnicity	height	income	\
0	working on college/university	asian	75.0	-1	
1	working on space camp	white	70.0	80000	
2	graduated from masters program	NaN	68.0	-1	
3	working on college/university	white	71.0	20000	
4	graduated from college/university	asian	66.0	-1	

	job	location	orientation	religion	sex	\
0	transportation	california	straight	agnosticism	m	
1	hospitality / travel	california	straight	agnosticism	m	
2	NaN	california	straight	NaN	m	
3	student	california	straight	NaN	m	
4	artistic / musical / writer	california	straight	NaN	m	


	sign	smokes	speaks	status
0	gemini	sometimes	english	single
1	cancer	no	english	single
2	pisces	no	english	available
3	pisces	no	english	single
4	aquarius	no	english	single

From the method `.head()` we can observe that the `income` has values equal to -1, which doesn't make sense for what the feature represents, meaning that these indicate missing values. To get an idea of the distribution of these values, we can plot a histogram of the feature:

```
In [6]: print('Missing income values: ', len(df[df.income == -1]))
print('Total values: ', len(df))
print('Percentage of missing income values: ', len(df[df.income == -1]) * 100 / len(df))

sns.displot(data=df, x="income", hue="sex", kind="hist", binwidth = 50000, multiple=True)
plt.show()
plt.clf()
```

```
Missing income values: 48442
Total values: 59946
Percentage of missing income values: 80.80939512227671 %
```



<Figure size 640x480 with 0 Axes>

Since more than 80% of the income values are missing, we will have to drop this feature as well.

```
In [7]: df = df.drop(columns=['income'])  
print(df.head())
```

	age	body_type	diet	drinks	drugs	\
0	22	a little extra	strictly anything	socially	never	
1	35	average	mostly other	often	sometimes	
2	38	thin	anything	socially	NaN	
3	23	thin	vegetarian	socially	NaN	
4	29	athletic	NaN	socially	never	

	education	ethnicity	height	\
0	working on college/university	asian	75.0	
1	working on space camp	white	70.0	
2	graduated from masters program	NaN	68.0	
3	working on college/university	white	71.0	
4	graduated from college/university	asian	66.0	


	job	location	orientation	religion	sex	\
0	transportation	california	straight	agnosticism	m	
1	hospitality / travel	california	straight	agnosticism	m	
2	NaN	california	straight	NaN	m	
3	student	california	straight	NaN	m	
4	artistic / musical / writer	california	straight	NaN	m	

	sign	smokes	speaks	status
0	gemini	sometimes	english	single
1	cancer	no	english	single
2	pisces	no	english	available
3	pisces	no	english	single
4	aquarius	no	english	single

We will now inspect our other numerical variables in order to make sure we don't have any more missing values or outliers.

```
In [8]: sns.displot(data=df, x="age", hue="sex", kind="hist", binwidth = 5, multiple = "stack")
plt.show()
plt.clf()

print('People that claim to be over 100: ', len(df[df.age > 100]))
print('People that claim to be over 68: ', len(df[df.age > 68]))
```




People that claim to be over 100: 2
 People that claim to be over 68: 33
 <Figure size 640x480 with 0 Axes>

Since in the feature `age` we have to people that claim to be over 100 years old, we can attribute this to be a mistake, therefore we will drop these values. In total, there are 33 people that claim to be over the age of 68, while this could be true, in order to avoid skewness in our data, all of these values will be dropped.

```
In [9]: df = df[df.age <= 68]
```

```
In [10]: sns.displot(data=df, x="height", hue="sex", kind="hist", binwidth = 5, multiple = "s")
plt.show()
plt.clf()

print('People tha claim to be over 80 in: ', len(df[df.height > 80]))
print('People tha claim to be under 55 in: ', len(df[df.height < 55]))
```





People tha claim to be over 80 in: 77
 People tha claim to be under 55 in: 39
 <Figure size 640x480 with 0 Axes>



In the feature `height` we have 77 people that claim to be over 80 in (203.2 cm) and 39 people that claim to be under 55 in (139.7 cm), while this could be true, omce again, we will eliminate these values to avoid skewness.



```
In [11]: df = df[(df.height >= 55) & (df.height <= 80)]
```


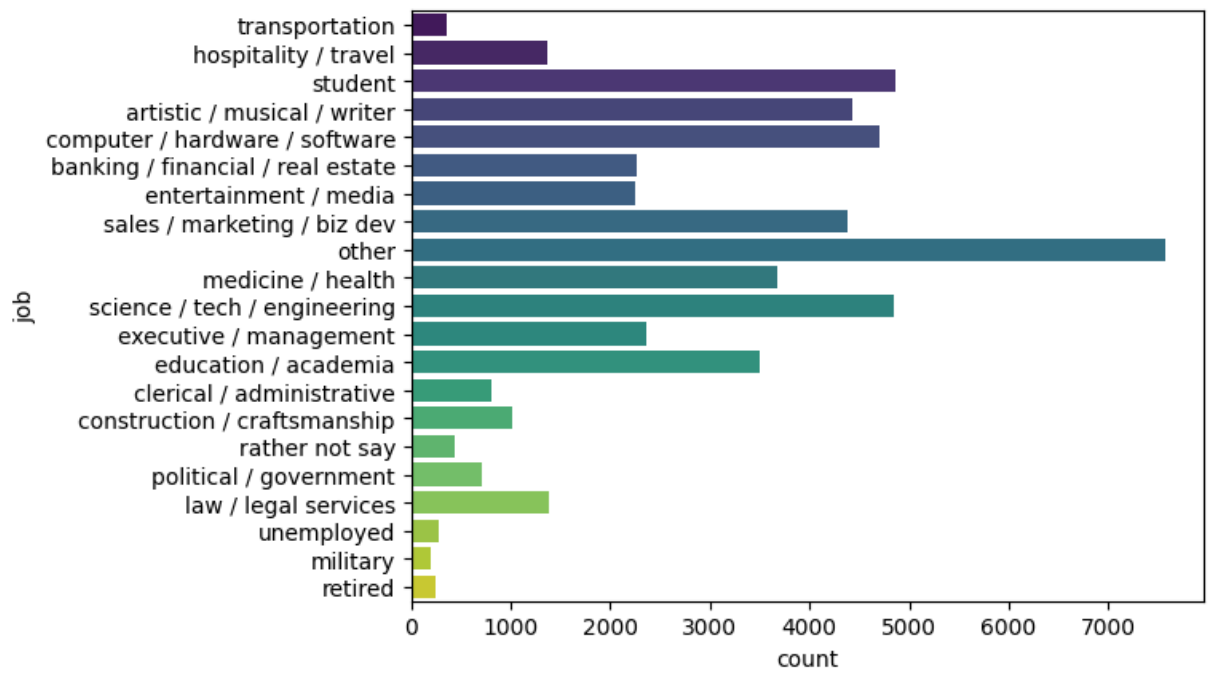
Now, we will plot the rest of our variables using a for loop.



```
In [12]: for column in df.columns:
          if column not in ['age', 'height']:
              sns.countplot(data=df, y=column, hue=column, legend=False, palette=sns.color_palette('magma'))
              plt.show()
              plt.clf()
```

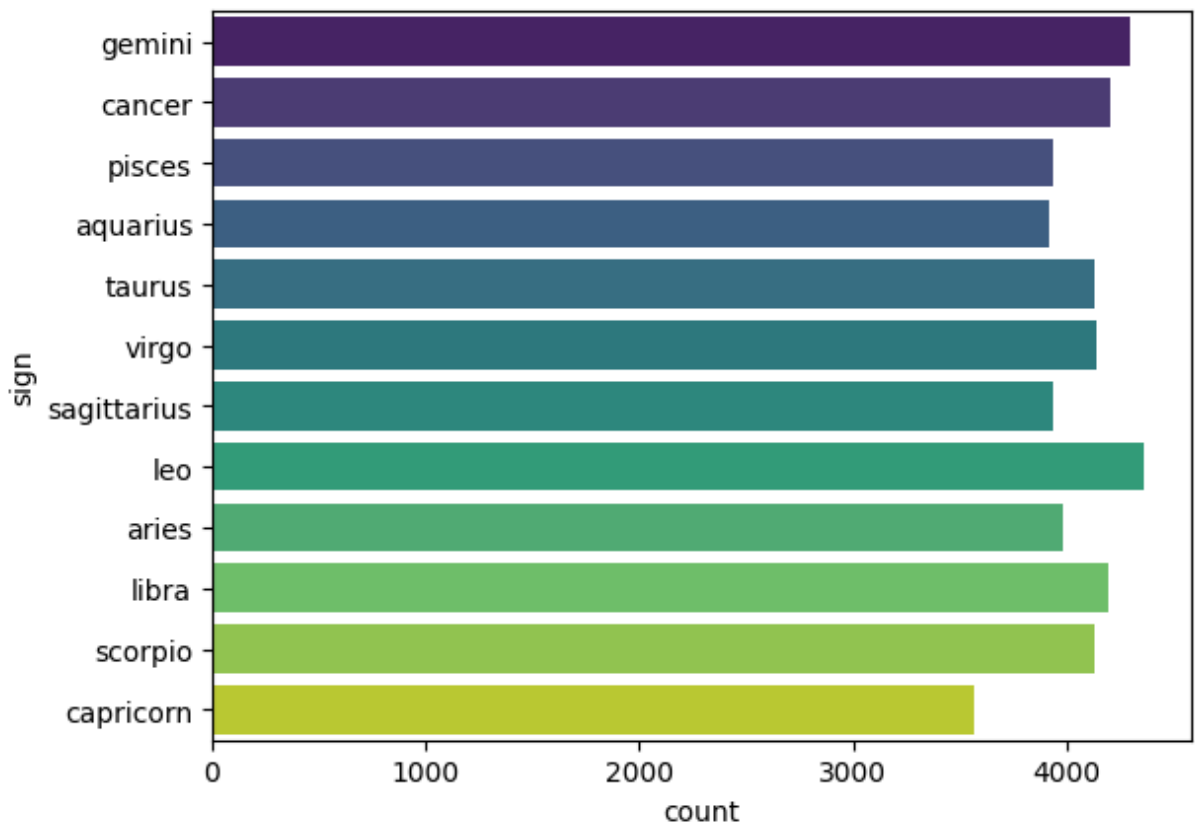




















<Figure size 640x480 with 0 Axes>

From these plots we can observe that:

1. Most people are from California, given that most of there's so little data about other locations, this feature will be removed.
2. Most people speak English as their primary language, so this feature will be removed for not having enough data about the rest as well.
3. Some people have their status as unknown, this is a clear mistake. While the most probable option is that they are also single, given that is such a short amount of data, these will be removed.
4. The status married is so rare that could affect the performance of our models later on, so these values will be ignored.
5. The values of the feature smokes could be grouped so we only have yes and no.
6. There are several other features that can be further simplified.

```
In [13]: df = df.drop(columns=['location', 'speaks'])
df = df[(df.status != 'unknown') & (df.status != 'married')]
df['smokes'] = df['smokes'].map(lambda x: 'yes' if x != 'no' else 'no')
```

```
In [14]: diet_mapping = {
    'mostly anything': 'anything', 'anything': 'anything', 'strictly anything': 'an
    'mostly vegetarian': 'vegetarian', 'vegetarian': 'vegetarian', 'strictly vegeta
    'mostly vegan': 'vegan', 'vegan': 'vegan', 'strictly vegan': 'vegan',
    'mostly other': 'other', 'other': 'other', 'strictly other': 'other',
    'mostly kosher': 'kosher', 'kosher': 'kosher', 'strictly kosher': 'kosher',
    'mostly halal': 'halal', 'halal': 'halal', 'strictly halal': 'halal'
}
```

```
df['diet'] = df['diet'].map(diet_mapping)
```

```
In [15]: body_type_mapping = {
    'thin': 'slim', 'skinny': 'slim',
    'average': 'average',
    'fit': 'athletic', 'athletic': 'athletic', 'jacked': 'athletic',
    'curvy': 'curvy_extra', 'a little extra': 'curvy_extra', 'full figured': 'curvy',
    'overweight': 'overweight',
    'used up': 'other', 'rather not say': 'other'
}

df['body_type'] = df['body_type'].map(body_type_mapping)
```

```
In [16]: education_mapping = {
    'high school': 'high_school', 'working on high school': 'high_school', 'dropped
    'graduated from two-year college': 'two_year_college', 'working on two-year col
    'dropped out of two-year college': 'two_year_college', 'two-year college': 'two
    'graduated from college/university': 'college_university', 'working on college/
    'dropped out of college/university': 'college_university', 'college/university'
    'graduated from masters program': 'masters', 'working on masters program': 'mas
    'dropped out of masters program': 'masters', 'masters program': 'masters',
    'graduated from ph.d program': 'phd_professional', 'working on ph.d program': '
    'dropped out of ph.d program': 'phd_professional', 'ph.d program': 'phd_profess
    'graduated from law school': 'phd_professional', 'working on law school': 'phd_
    'dropped out of law school': 'phd_professional', 'law school': 'phd_professiona
    'graduated from med school': 'phd_professional', 'working on med school': 'phd_
    'dropped out of med school': 'phd_professional', 'med school': 'phd_professiona
    'graduated from space camp': 'space_camp', 'working on space camp': 'space_camp
    'dropped out of space camp': 'space_camp', 'space camp': 'space_camp'
}

df['education'] = df['education'].map(education_mapping)
```

```
In [17]: job_mapping = {
    'computer / hardware / software': 'stem', 'science / tech / engineering': 'stem
    'artistic / musical / writer': 'creative_media', 'entertainment / media': 'crea
    'sales / marketing / biz dev': 'business_finance', 'banking / financial / real
    'executive / management': 'business_finance',
    'medicine / health': 'healthcare_law', 'law / legal services': 'healthcare_law'
    'education / academia': 'education_government', 'political / government': 'educ
    'hospitality / travel': 'trades_services', 'construction / craftsmanship': 'tra
    'clerical / administrative': 'trades_services', 'transportation': 'trades_servi
    'unemployed': 'unemployed_retired', 'retired': 'unemployed_retired',
    'rather not say': 'other_unknown', 'other': 'other_unknown',
    'military': 'military',
    'student': 'student'
}

df['job'] = df['job'].map(job_mapping)
```

Now we can drop the rest of null values to have a completely clean dataset.

```
In [18]: new_df = df.dropna()
print(new_df.head())
```

	age	body_type	diet	drinks	drugs	education	\
0	22	curvy_extra	anything	socially	never	college_university	
1	35	average	other	often	sometimes	space_camp	
7	31	average	anything	socially	never	college_university	
9	37	athletic	anything	not at all	never	two_year_college	
11	28	average	anything	socially	never	college_university	

	ethnicity	height	job	orientation	religion	sex	\
0	asian	75.0	trades_services	straight	agnosticism	m	
1	white	70.0	trades_services	straight	agnosticism	m	
7	white	65.0	creative_media	straight	christianity	f	
9	white	65.0	student	straight	atheism	m	
11	white	72.0	business_finance	straight	christianity	m	

	sign	smokes	status
0	gemini	yes	single
1	cancer	no	single
7	sagittarius	no	single
9	cancer	no	single
11	leo	no	seeing someone

We can also show some summary statistics for our data:

```
In [19]: print(new_df.describe())
for column in new_df.columns:
    if column not in ['age', 'height']:
        print(f'\nValue counts for {column}: ')
        print(new_df[column].value_counts())
print(new_df.shape)
```

	age	height
count	14478.000000	14478.000000
mean	32.788023	68.248791
std	10.144015	3.846002
min	18.000000	55.000000
25%	26.000000	65.000000
50%	30.000000	68.000000
75%	38.000000	71.000000
max	68.000000	80.000000

Value counts for body_type:

body_type	
athletic	6266
average	3907
curvy_extra	2349
slim	1669
other	149
overweight	138

Name: count, dtype: int64

Value counts for diet:

diet	
anything	11361
vegetarian	1952
other	798
vegan	276
kosher	60
halal	31

Name: count, dtype: int64

Value counts for drinks:

drinks	
socially	10275
rarely	1669
often	1252
not at all	1049
very often	142
desperately	91

Name: count, dtype: int64

Value counts for drugs:

drugs	
never	11504
sometimes	2825
often	149

Name: count, dtype: int64

Value counts for education:

education	
college_university	8661
masters	2924
phd_professional	1150
two_year_college	1107
space_camp	533
high_school	103

Name: count, dtype: int64

Value counts for ethnicity:

ethnicity

white	8908
asian	2142
hispanic	1173
black	850
other	433
indian	361
middle	218
pacific	198
native	195

Name: count, dtype: int64

Value counts for job:

job

stem	2607
business_finance	2498
other_unknown	2139
creative_media	1836
student	1532
healthcare_law	1372
education_government	1226
trades_services	1007
unemployed_retired	198
military	63

Name: count, dtype: int64

Value counts for orientation:

orientation

straight	12622
gay	1198
bisexual	658

Name: count, dtype: int64

Value counts for religion:

religion

agnosticism	3127
other	2939
atheism	2318
christianity	2233
catholicism	1881
judaism	970
buddhism	748
hinduism	209
islam	53

Name: count, dtype: int64

Value counts for sex:

sex

m	8603
f	5875

Name: count, dtype: int64

Value counts for sign:

sign

```

gemini      1314
cancer      1302
virgo       1277
leo         1258
libra       1242
scorpio     1207
taurus      1200
aries       1190
pisces      1165
sagittarius 1155
aquarius    1112
capricorn   1056
Name: count, dtype: int64

```

```

Value counts for smokes:
smokes
no      11516
yes      2962
Name: count, dtype: int64

```

```

Value counts for status:
status
single      13642
seeing someone  428
available    408
Name: count, dtype: int64
(14478, 15)

```

2.3. Data Preprocessing

Now that we have a clean dataset, we will have to transform it's values into new ones that our models will be able to accept. This will be done through the methods of One-Hot encoding and Standardization.

```

In [20]: columns = ['age', 'body_type', 'diet', 'drinks', 'smokes', 'drugs', 'education', 'j
data = new_df[columns]

```

```

In [21]: from sklearn.preprocessing import StandardScaler

num_features = ['age', 'height']
cat_features = ['body_type', 'diet', 'drinks', 'drugs', 'education', 'job', 'religi

data.loc[:, 'sex'] = data['sex'].map({'m': 0, 'f': 1})
data.loc[:, 'smokes'] = data['smokes'].map({'no': 0, 'yes': 1})

scaler = StandardScaler()
for num in num_features:
    data.loc[:, num] = scaler.fit_transform(data[[num]])

cat_data = pd.get_dummies(data[cat_features], drop_first=True).astype(int)
data_preprocessed = pd.concat([data[num_features + ['sex', 'smokes']], cat_data], a

print(data_preprocessed.head())

```

	age	height	sex	smokes	body_type_average	body_type_curvy_extra	\
0	-1.063523	1.755444	0	1	0	1	
1	0.218065	0.455348	0	0	1	0	
7	-0.176270	-0.844748	1	0	1	0	
9	0.415232	-0.844748	0	0	0	0	
11	-0.472021	0.975387	0	0	1	0	

	body_type_other	body_type_overweight	body_type_slim	diet_halal	...	\
0	0	0	0	0	0	...
1	0	0	0	0	0	...
7	0	0	0	0	0	...
9	0	0	0	0	0	...
11	0	0	0	0	0	...

	religion_atheism	religion_buddhism	religion_catholicism	\
0	0	0	0	
1	0	0	0	
7	0	0	0	
9	1	0	0	
11	0	0	0	

	religion_christianity	religion_hinduism	religion_islam	\
0	0	0	0	
1	0	0	0	
7	1	0	0	
9	0	0	0	
11	1	0	0	

	religion_judaism	religion_other	orientation_gay	orientation_straight
0	0	0	0	1
1	0	0	0	1
7	0	0	0	1
9	0	0	0	1
11	0	0	0	1

[5 rows x 45 columns]

C:\Users\sergi\AppData\Local\Temp\ipykernel_8204\4048868718.py:11: FutureWarning: Setting an item of incompatible dtype is deprecated and will raise in a future error of pandas. Value '[-1.06352322 0.21806485 -0.17626994 ... -0.86635582 0.90815074 -0.57060473]' has dtype incompatible with int64, please explicitly cast to a compatible dtype first.

```
data.loc[:, num] = scaler.fit_transform(data[[num]])
```


We can now assign values for our independent variable (X) and our target variable (y):

```
In [22]: X = data_preprocessed.drop(columns='sex')
y = data_preprocessed['sex']
y = y.astype(int)
```

It's always a good practice to look for multicollinearity in our data, therefore we plot the correlation matrix using a heatmap:

```
In [23]: corr_matrix = X.corr(method='pearson')
sns.heatmap(corr_matrix, cmap='RdBu_r')
```

```
plt.show()
plt.clf()
```



<Figure size 640x480 with 0 Axes>

Now that we've made sure that our data is ready for modelling, we finally can split our data for proper training and testing:

```
In [24]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.8, test_size=0.2)
```

3. Modelling

3.1. Logistic Regression

In this section we will implement an optimized version of logistic regression using BayesSearchCV.

```
In [25]: from skopt import BayesSearchCV
from sklearn.linear_model import LogisticRegression

lr_opt = BayesSearchCV(
```

```

LogisticRegression(solver='liblinear', random_state=1),
{
    'C': (1e-6, 1e+6, 'log-uniform'),
    'penalty': ['l1', 'l2'],
},
n_iter=32,
cv=5
)

lr_opt.fit(X_train, y_train)

print("Best parameters:", lr_opt.best_params_)
print("Validation score:", lr_opt.best_score_)
print("Test score:", lr_opt.score(X_test, y_test))

```

Best parameters: OrderedDict({'C': 0.7010828232480305, 'penalty': 'l1'})
Validation score: 0.8647896862046167
Test score: 0.8756906077348067

Now we print the classification report to take a look at our model's scores:

In [26]: `from sklearn.metrics import classification_report`

```

y_pred = lr_opt.predict(X_test)
print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
0	0.88	0.91	0.90	1704
1	0.87	0.82	0.84	1192
accuracy			0.88	2896
macro avg	0.87	0.87	0.87	2896
weighted avg	0.88	0.88	0.88	2896

Finally, we plot the confusion matrix using a heatmap:


In [27]: `from sklearn.metrics import confusion_matrix`

```

def plot_cm(y_test, y_pred):
    cart_cm = confusion_matrix(y_test, y_pred)
    ax = plt.subplot()
    sns.heatmap(cart_cm, annot=True, ax = ax, fmt="d");
    ax.xaxis.set_ticklabels(['Male', 'Female']);
    ax.yaxis.set_ticklabels(['Male', 'Female']);
    plt.show()
    plt.clf()

plot_cm(y_test, y_pred)

```



<Figure size 640x480 with 0 Axes>

3.2. Decision Tree

For this section, we will use a decision tree instead to see if we can achieve better results. Like in the previous section we will first create and fit the model, then we will print the classification report and lastly we will plot the confusion matrix.

```
In [28]: from sklearn.tree import DecisionTreeClassifier

dt_opt = BayesSearchCV(
    DecisionTreeClassifier(random_state=1),
    {
        'max_depth': (1, 50),
        'min_samples_split': (2, 20),
        'min_samples_leaf': (1, 20),
        'max_features': ['sqrt', 'log2', None],
        'criterion': ['gini', 'entropy']
    },
    n_iter=32,
    cv=5
)

dt_opt.fit(X_train, y_train)


print("Best parameters:", dt_opt.best_params_)
print("Validation score:", dt_opt.best_score_)
print("Test score:", dt_opt.score(X_test, y_test))
```

Best parameters: OrderedDict({'criterion': 'gini', 'max_depth': 45, 'max_features': None, 'min_samples_leaf': 20, 'min_samples_split': 20})
 Validation score: 0.8484713870520736
 Test score: 0.8515193370165746

```
In [29]: y_pred = dt_opt.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.90	0.88	1704
1	0.85	0.78	0.81	1192
accuracy			0.85	2896
macro avg	0.85	0.84	0.84	2896
weighted avg	0.85	0.85	0.85	2896

```
In [30]: plot_cm(y_test, y_pred)
```



<Figure size 640x480 with 0 Axes>

3.3. Random Forest

For this section, we will use a random forest classifier instead to see if we can achieve better results. Like in the previous section we will first create and fit the model, then we will print the classification report and lastly we will plot the confusion matrix.

```
In [31]: from sklearn.ensemble import RandomForestClassifier
```

```

rfc_opt = BayesSearchCV(
    RandomForestClassifier(n_jobs=-1, random_state=1),
    {
        'n_estimators': (50, 500),
        'max_depth': (5, 50),
        'min_samples_split': (2, 20),
        'min_samples_leaf': (1, 20),
        'max_features': ['sqrt', 'log2', None],
        'bootstrap': [True, False],
        'criterion': ['gini', 'entropy']
    },
    n_iter=32,
    cv=5
)

rfc_opt.fit(X_train, y_train)

print("Best parameters:", rfc_opt.best_params_)
print("Validation score:", rfc_opt.best_score_)
print("Test score:", rfc_opt.score(X_test, y_test))

```

Best parameters: OrderedDict({'bootstrap': True, 'criterion': 'gini', 'max_depth': 50, 'max_features': 'sqrt', 'min_samples_leaf': 5, 'min_samples_split': 2, 'n_estimators': 376})

Validation score: 0.859955029395256

Test score: 0.8705110497237569

```

In [32]: y_pred = rfc_opt.predict(X_test)
         print(classification_report(y_test, y_pred))


```

	precision	recall	f1-score	support
0	0.87	0.92	0.89	1704
1	0.87	0.81	0.84	1192
accuracy			0.87	2896
macro avg	0.87	0.86	0.86	2896
weighted avg	0.87	0.87	0.87	2896

```

In [33]: plot_cm(y_test, y_pred)

```

<Figure size 640x480 with 0 Axes>

3.4. Support Vector Machine

In [34]: `from sklearn.svm import SVC`

```
svm_opt = BayesSearchCV(
    SVC(kernel="rbf"),
    {
        'C': (1e-3, 1e+2, 'log-uniform'),
        'gamma': (1e-4, 1, 'log-uniform'),
    },
    n_iter=20,
    cv=3
)

svm_opt.fit(X_train, y_train)


print("Best parameters:", svm_opt.best_params_)
print("Validation score:", svm_opt.best_score_)
print("Test score:", svm_opt.score(X_test, y_test))
```

Best parameters: OrderedDict({'C': 0.9662810507409029, 'gamma': 0.1309846024920814})
 Validation score: 0.8655669444097769
 Test score: 0.8756906077348067

In [35]: `y_pred = rfc_opt.predict(X_test)`
`print(classification_report(y_test, y_pred))`

	precision	recall	f1-score	support
0	0.87	0.92	0.89	1704
1	0.87	0.81	0.84	1192
accuracy			0.87	2896
macro avg	0.87	0.86	0.86	2896
weighted avg	0.87	0.87	0.87	2896

In [36]: `plot_cm(y_test, y_pred)`



<Figure size 640x480 with 0 Axes>

4. Conclusion

In this analysis, we aimed to predict users' gender on the OKCupid dataset using four machine learning models: **Logistic Regression**, **Decision Tree**, **Random Forest**, and **Support Vector Machine (SVM)**. After preprocessing the data by encoding categorical variables, standardizing numerical features, and performing hyperparameter tuning, we evaluated the models based on their validation and test scores.

- **Best Performing Model:** The **SVM** model achieved the highest test score of **0.8809**, indicating its strong ability to generalize.
- **Efficiency vs. Performance Tradeoff:** While **Random Forest** and **SVM** provided competitive results, they took significantly longer to train compared to **Logistic Regression**, which performed nearly as well in just 21 seconds.

- **Decision Tree Weakness:** The **Decision Tree** had the lowest test score (**0.8567**), likely due to its tendency to overfit.

Overall, **SVM** provided the best results but at the cost of processing time, while **Logistic Regression** was a strong contender with much lower computational demand.

4.1. Next Steps for Further Analysis

1. Feature Engineering & Selection

- Explore new features that could enhance prediction accuracy, such as **profile text analysis (NLP techniques)**.

2. Ensemble Learning

- Combine multiple models (e.g., **Stacking or Boosting**) to leverage their strengths and improve prediction robustness.

3. Deep Learning Approach

- Implement a **Neural Network (MLP)** and compare its performance with traditional models.

4. Explainability & Bias Analysis

- Investigate whether certain features disproportionately affect gender classification, ensuring the model is fair and interpretable.